

struggle with

Q7.1 — "No field agents, no home visits. Then who shows a 52-year-old weaver how to open your app the first time?"

Think. This is the adoption question and our answer is a constraint, not a plan — so be honest that it is a bet, and name what the bet rests on.

The constraint, and it is deliberate (docs/Master-Technical-Reference.md §2, three hard NOs): **no field agents, no human visits** — because a solution requiring a visit breaks the PS's own words, "*without requiring advanced technical knowledge*." If the app only works when someone sits beside them, we have built a training programme with a login screen.

So the first run has to carry itself, and that is what the eight-screen onboarding is for: language tiles that **speak their own name on tap**, a spoken consent notice, and **the camera by screen eight** — no tour, no dashboard, no feature list. The very first thing that happens is the thing they came for.

And the camera is the part that genuinely needs no teaching. Red ring, shutter greyed out and dead. Green ring, shutter glows. Hold green one second, it fires itself. **They move until the button lights up.** No instruction, no literacy, no helper.

What exists behind it when that is not enough:

- **/help — one button, one callback request to a cluster coordinator.** Remote, a phone call, never a visit.
- **The admin console** is built for exactly this: assisted onboarding support, remote (docs/Application-Architecture.md §9).
- **Existing human infrastructure we did not have to build:** ~5.4 lakh **CSC Village Level Entrepreneurs**, ~4.35 lakh of them in gram panchayats, plus cluster coordinators at Common Facility Centres. **These people already exist, are already funded, and already sit where our users are.**

Breaks when. The bet is that a spoken, gated, three-tap flow clears the first-run barrier without a person present. **We have not tested it with a single artisan.** That is the largest untested assumption in the project and it is not a technical one. One afternoon in one cluster would tell us more than another week of code.

Q7.2 — "Your impact goal is higher income. How would you prove it, and what if you are wrong?"

Think. Give the measurement design, then state the falsification condition. A team that names what would prove it wrong is more credible than one that only names what would prove it right.

What we can measure honestly:

Signal	Source	Honest?
Listings created per artisan	Ours	Yes — but it is an output, not an outcome
Orders per artisan per channel	Order	Yes. The real number
Realised price vs the computed floor	<code>floor_price</code> vs price	Yes — directly measures whether we stopped a loss-making sale
artisan_confirmed_payment	The "Paisa aaya?" tap	Yes — settlement reality, not the marketplace's promise
Total annual income	—	No. We cannot see their bank account

So the defensible claim is narrower than the PS's phrasing and we should use the narrow one: we can prove **channel access created, orders received, and sales prevented below cost**. We cannot prove total household income, and any team claiming they can is not counting the middleman sales that never touch a platform.

The falsification condition, stated up front: if listings per artisan rise and **orders per artisan stay flat**, then cataloging was not the bottleneck and our thesis is wrong. That number is in the schema from day one specifically so it can embarrass us.

The counterfactual problem, conceded: an artisan who adopts our app is already more motivated than one who does not. Any income difference is partly selection. **The only clean answer is a staggered rollout across comparable clusters** — cluster A onboards in month 1, cluster B in month 4 — which a ministry pilot can actually do and a hackathon cannot.

Q7.3 — "Who pays for this? Your slide says government support, white-label licensing and marketplace. Which one is real?"

Think. Answer the constraint first, because it is the part that binds.

The hard constraint (`docs/Master-Technical-Reference.md §2`): **we never charge artisans. Anything.** That is a locked no, and it removes the obvious revenue line.

Which leaves, in order of realism for a government deployment:

1. **Government-funded, per-artisan-onboarded.** The natural fit for MoSJE / DC-Handicrafts / PM Vishwakarma, and it is the only model where incentives align — the ministry wants artisans digitised, we are paid to digitise them. **Costing document gives the per-artisan figure**, and it is small enough that this is a rounding error against existing scheme outlays.

2. **White-label to state handicraft boards and cooperatives.** Odisha, Gujarat, UP each run their own emporium and cluster programmes. Same code, per-tenant, and multi-tenancy is already in the design.
3. **Marketplace commission on our own surface only** — never on GeM or ONDC, where the artisan's sale is theirs.

The uncomfortable honest note: if the government funds it and then stops, the running cost falls on whoever holds the servers. That is why the costing document's most important number is not the total — it is **cost per artisan per year**, and why every cost-reduction lever in it matters.

Q7.4 — "Environmental benefit through reduced travel. Is that a real claim or a slide filler?"

Answer. Directionally true, unquantified, and I would not lead with it.

The mechanism is real: an artisan who sells year-round from a phone travels less to periodic fairs and melas, and **the Cluster Hub consolidates fifty rural courier pickups into one stop** — that second one is a genuine, countable reduction in vehicle-kilometres.

But we have measured nothing, and there is an offsetting cost we do not currently count: parcel shipping for individual online orders, which fair-based selling does not incur. **A claim that ignores its own counter-term is not a claim.**

Keep it on the slide, phrase it as the hub consolidation — that one survives scrutiny. Drop the general travel-emissions framing.

Q7.5 — "Return to origin. A buyer refuses delivery. Who eats the cost?"

Think. This is the question that quietly decides whether the artisan is better or worse off, and it deserves a straight answer.

The decision: prepaid only at launch. COD is huge in India and carries huge RTO risk, and **artisans cannot absorb return-to-origin costs.** One refused COD parcel can erase the margin on five sales.

And the rule that follows: return-shipping liability is **defined explicitly per channel and shown to the artisan BEFORE they list** — **never a surprise.** Also enforced before listing: **pincode serviceability.** If their pincode is not serviceable, we block self-ship channels and route to the Cluster Hub only. **Discovering that after a sale is a disaster**, and it is why pincode is collected on onboarding screen 6 rather than at checkout.

Who inspects a returned item? The Cluster Hub. **We track RTO rate per artisan and per category**, because it is the metric that predicts whether a category is viable online at all.

Q7.6 — "You have four weeks. What are you cutting?"

Think. Show the build order, and show that the cut line was drawn before the pressure arrived.

Phases 1–6 are the defensible core, and each phase ends in something demoable with nothing built that a later phase rewrites:

Phase	Ships	Proves
1	Onboarding, camera gate, resumable upload, schema	An artisan can sign up and take a good photo
2	ai/enhance real + colour lock	The main image is marketplace-grade
3	Bhashini ASR/TTS + ai/catalog	PS feature 2 done
4	ai/price cost-up + floor + GeM discount math	PS feature 3 done
5	Marketplace SSR + /publish	First real one-tap listing
6	ONDC MSN on staging + adapter	Tier A complete. The USP is real

Phases 11 and 12 — the guided browser with autofill, and generative secondary images — are upside. If either slips, nothing in the PS goes unanswered. That sentence is in docs/Application-Architecture.md §11, written before the deadline pressure, which is the only time a cut line means anything.

And the generative feature specifically is scoped as what it is: "a demo weapon, not a product foundation. 20 seconds on stage. Build the boring pipeline properly."

The five questions we would struggle with

Written down so nobody is surprised on stage. **Each one has a real answer that is a concession.**

1. "Have you put this in front of a single artisan?"

No. Every threshold is calibrated on **web-sourced** fixtures, not photographs taken by artisans on ₹7,000 phones in Indian light. Our own research/RESULTS.md says the fixture set is "too web-sourced." **This is the honest ceiling on every number in Part 2**, and no amount of further engineering removes it.

2. "Your slide says ML pricing and your code has none."

Concede immediately, explain that arithmetic is the defensible choice with no training data, and **fix the slide before the presentation** (Part 0, A0.1). If the panel finds it first, everything else we say gets re-read as marketing.

3. "Show me a live listing on GeM or Amazon."

We cannot. GeM has no API and we have **zero real category templates**. Amazon's SP-API PUT is unimplemented, public-app publication requires Amazon review that is not achievable in this timeframe, and public apps are capped at 25 seller authorisations until then. **We can show our own marketplace end to end**, and we should offer that proactively rather than be walked into the gap.

4. "Prove your app increases income."

We can prove orders, channel access, and losses prevented at the floor. **We cannot prove income** — we cannot see their bank account, and middleman sales never touch a platform. Only a staggered ministry pilot answers this properly.

5. "What if listing was never the bottleneck?"

The honest answer is that IndiaHandmade's 0.1% is **consistent** with our thesis and does not **prove** it. If listings rise and orders stay flat, we automated the wrong half. We built the metric that would tell us, and we would rather find out than not.

The pattern in all five: we know exactly where the edge is. A team that can point at its own edge is more trustworthy than one that claims not to have one — and on a government problem statement, trustworthiness is the thing being procured.